

Rowan Arora

510-468-2771 | rowan.arora@icloud.com | [Portfolio](#) | [LinkedIn](#) | [GitHub](#)

ABOUT ME

Recent computer science graduate with expertise in **fullstack development** and **scalable architecture**. Skilled in **Python**, **FastAPI**, **MongoDB**, and cloud services, with hands-on experience building production-ready APIs and integrating **AI/LLM APIs** (OpenAI GPT-4, Grok, Google Gemini) for document and contract workflows. Proven ability to collaborate with senior engineers in **Agile environments** and deliver high-impact solutions in fast-paced tech companies.

EDUCATION

University of Toronto, St. George
Honors Bachelor of Science, Computer Science

Toronto, ON
Graduated: June 2025

TECHNICAL SKILLS AND CERTIFICATIONS

Languages: Python, Java, SQL, C++, C, HTML, CSS, JavaScript, TypeScript.

Frameworks & APIs: FastAPI, REST APIs, Pydantic, SQLAlchemy, Vue.js, React, Django, Flask.

Cloud & Databases: MongoDB, PostgreSQL, Azure Service Bus, Azure Functions.

Tools & Platforms: Git, Bitbucket, Jira, Docker, Postman, pytest, pandas, Vite, LaTeX, scikit-learn, nltk, BeautifulSoup.

Certifications: Coursera Deep Learning Specialization (Neural Networks, CNNs, Sequence Models, etc.).

EXPERIENCE

Junior Python Developer | *Python, FastAPI, MongoDB, API Development, Azure, Vue.js* May 2025 – Present
Leah (formerly ContractPodAI) Toronto, ON

- **Developed and debugged backend APIs** using **FastAPI** for Leah, a proprietary AI assistant for legal contract redlining, building **15+ new endpoints** from scratch and fixing existing functionality to serve frontend and database integration needs, improving API reliability and response times.
- **Engineered document ingestion pipeline** using **Azure Functions** and **Azure Service Bus**, optimizing data processing workflows and reducing pipeline overhead by **15%** through efficient document processing and integration with **OpenAI's GPT-4, Grok, and Google's Gemini** APIs.
- **Designed and implemented database schemas** and models in **MongoDB**, managing **50+ collections** for uploaded documents, internal response/request models, and product data, ensuring optimal data structure for the legal tech platform and improving data organization.
- Collaborated with **7+ senior software engineers** in an Agile environment, handling **2-3 tickets daily** via **Jira** and **Bitbucket**, participating in code reviews and contributing to production-ready features deployed to **1000+ active users**.
- **Addressed security vulnerabilities** from internal reports: implemented backend **input validation and sanitization** to prevent LLM responses from disclosing sensitive data, added frontend **input parsing** to mitigate injection attacks, and applied fixes across the stack to strengthen security for legal contract workflows.
- Contributed to frontend fixes and features using **Vue.js**, supporting fullstack delivery and improving collaboration with the frontend team.

Software/Robotics Intern | *Fusion 360, ROS 2, Ignition Gazebo, URDF, Docker, Ubuntu* June 2024 – August 2024
Evodyne Robotics Academy Mountain View, CA

- **Developed a virtual simulation environment** for the Evodog robotic model using **CAD modeling (Fusion 360)**, **ROS 2**, and **Ignition Gazebo**, enabling physics simulations that eliminated the need for physical assembly and 3D printing.
- **Created a workflow** that replaced manual testing with virtual simulations, reducing resource consumption by **50%** and saving **100+ hours** of testing time while improving design accuracy by **25%**.
- **Delivered a fully functional virtual model** within **5 weeks**, mastering **URDF files**, **Docker**, and **Ubuntu** to enhance robotics development and testing efficiency.

PROJECTS

JournalismJobForge | *Python, FastAPI, SQLAlchemy, PostgreSQL, Web Scraping, Vue.js, Vite* February 2025 – Present

- Building a vertical job-search platform for journalism and media roles with a **FastAPI + SQLAlchemy + PostgreSQL** backend, modeling users, applications, resumes, and cover letters along with rich preferences (locations, titles, role types, skills) and configurable salary normalization / location-to-currency inference.
- Engineered a **JournalismJobs.com** scraper using **requests**, **BeautifulSoup**, and **tqdm** that batch-crawls up to **500 listings per page**, scraping and normalizing **thousands** of jobs per run, deduplicates by application URL, and parses noisy HTML into strongly-typed salary, job-type, industry, and currency fields to power salary-aware filtering and ranking.
- Designed REST endpoints for querying and paginating jobs, enforcing user-specific filters (preferred locations, role types, salary range), and supporting saved searches and application tracking via strongly-typed Pydantic models and reusable parsing utilities.
- Architecting a self-hosted web frontend using **Vue.js**, **Vite**, and component-driven UI patterns to surface personalized job feeds, saved searches, and application timelines, with planned **AI-assisted** resume/cover letter tailoring from stored documents and scraped job descriptions.